CS50 Section

Week 7

Attendance Sign In: tinyurl.com/sqlcs50rw

Agenda

- Week 6: Python - Remaining Questions
- Week 7: SQL

Week 6 Review - What Questions do You Have?

- For Loops
  - range()
  - len()
- While Loops
- Tuples: ()
- Lists: []
- Dictionaries: {}
- Strings
  - Substring Slicing: s[i:j]

Week 7

MOVING BEYOND THE CSV FILE

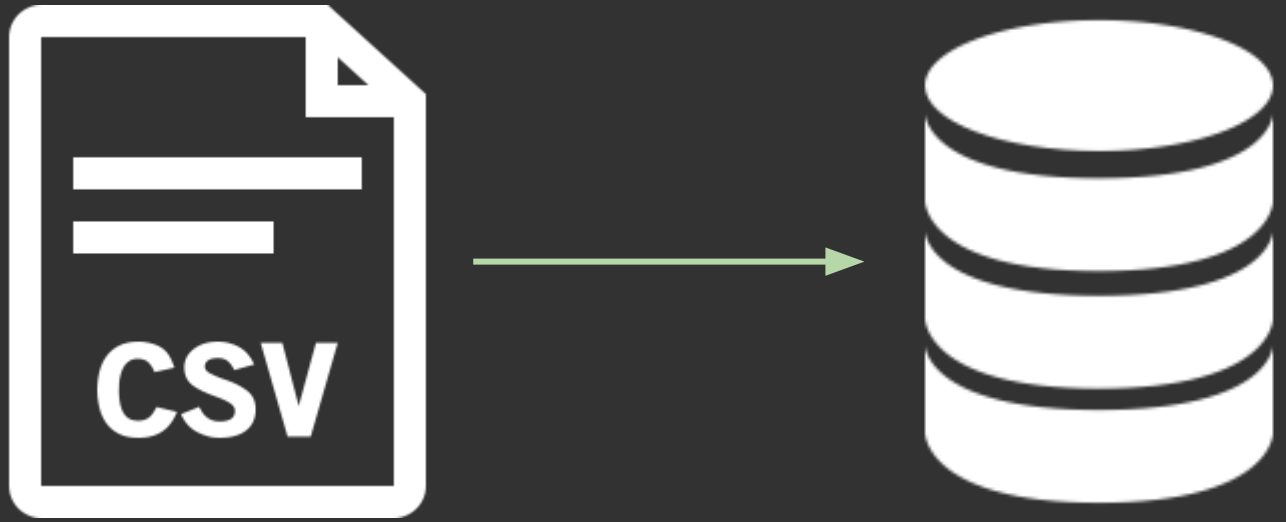

MOVING BEYOND THE CSV FILE

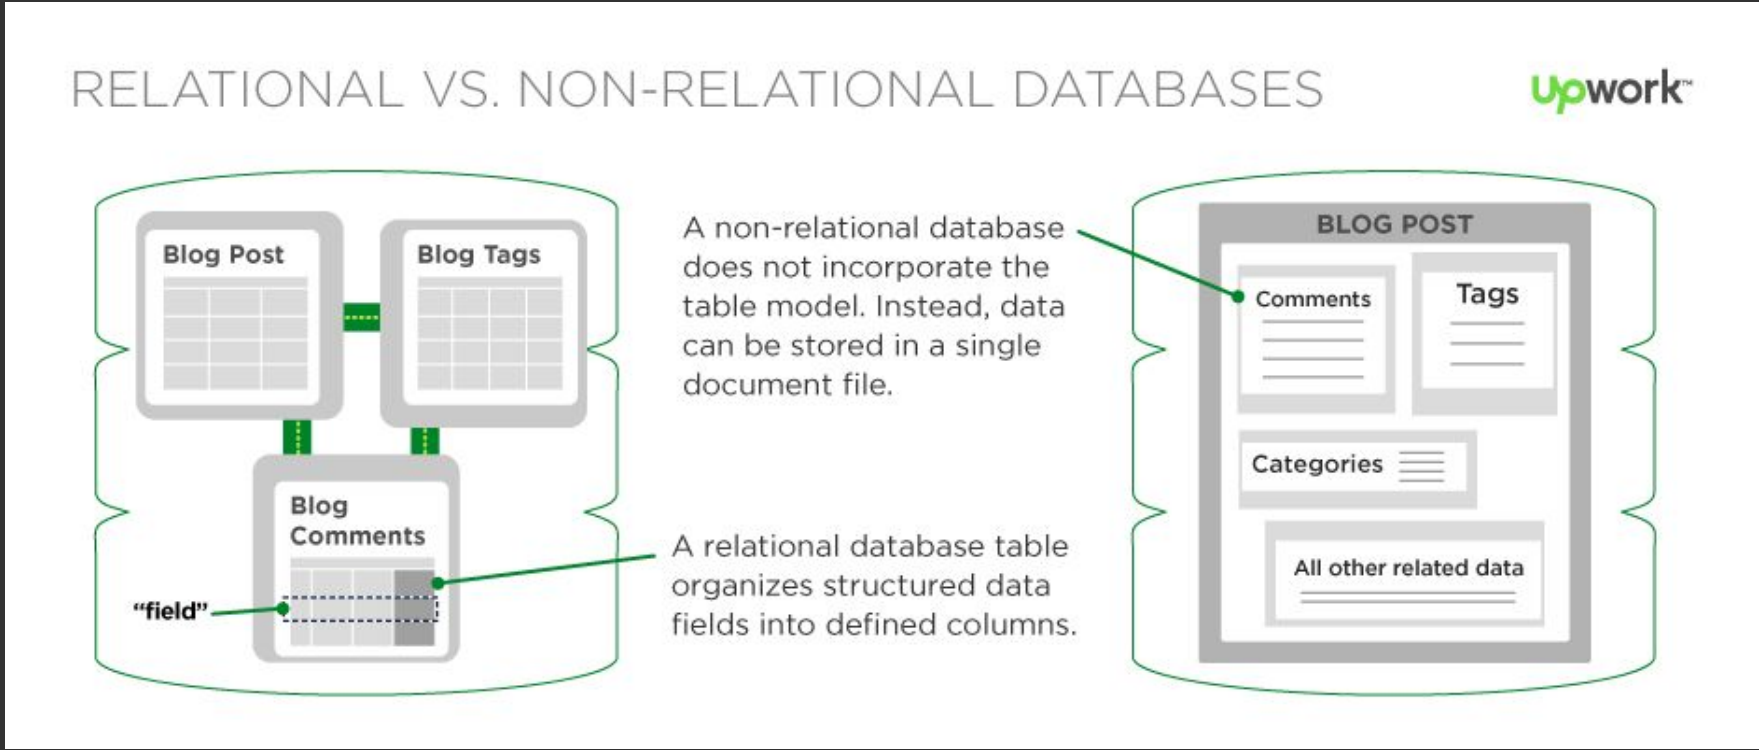

SQL

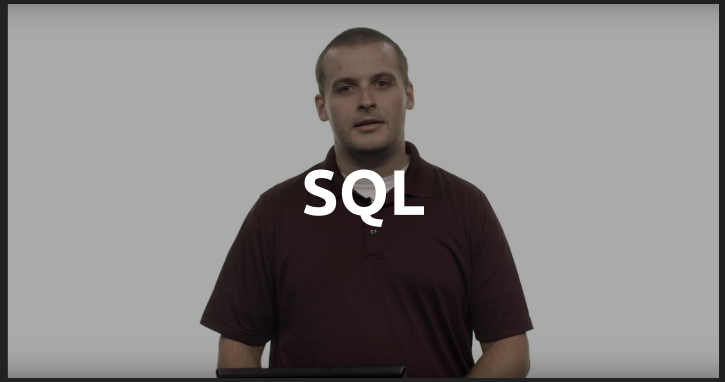

https://youtu.be/nfGiGSCEYRI

- As a relational DB, SQL is centered around **tables**.
- For each table, you specify all of the **columns** in the table.
- When new information is added to the database, the new information (typically) goes into a new **row**.
- There are many data types that can be stored in a SQL database. This is just a small sample.

INT	SMALLINT	TINYINT	MEDIUMINT	BIGINT
DECIMAL	FLOAT	BIT	DATE	TIME
DATETIME	TIMESTAMP	CHAR	VARCHAR	BINARY
BLOB	TEXT	ENUM	GEOMETRY	LINESTRING

- As a relational DB, SQL is centered around **tables**.
- For each table, you specify all of the **columns** in the table.
- When new information is added to the database, the new information (typically) goes into a new **row**.
- In SQLite, which we'll use in this course, we can consolidate these various datatypes into a few more general classes (though underlying types still exist)

NULL	INTEGER	REAL	TEXT	BLOB

DATABASE DESIGN PRINCIPLES

- Why do we use types when designing our databases?

DATABASE DESIGN PRINCIPLES

- Why do we use types when designing our databases?

They allow us to increase the efficiency of data storage/retrieval AND they improve data integrity.

DATABASE DESIGN PRINCIPLES

- Why do we use types when designing our databases?

A - Atomicity

All or Nothing Transactions

C - Consistency

Guarantees Committed Transaction State

I - Isolation

Transactions are Independent

D – Durability

Committed Data is Never Lost

(c) http://blog.sqlauthority.com

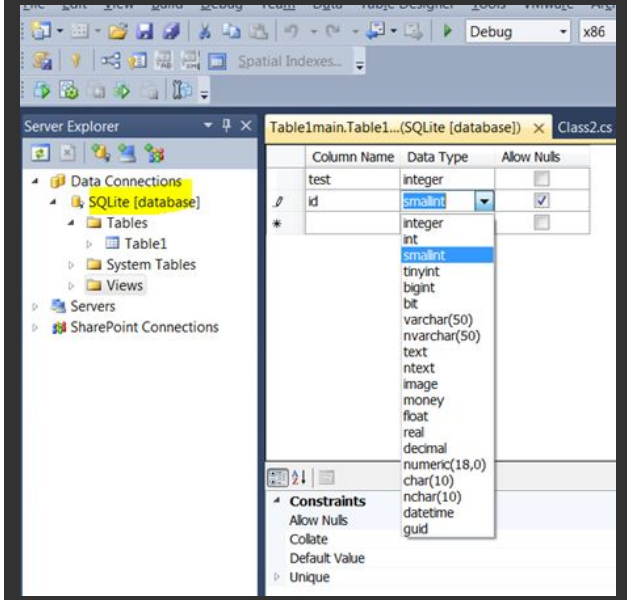

RELATIONAL DATABASE TERMINOLOGY

- **Database -** One discrete data system; The largest structure we'll be working with

RELATIONAL DATABASE TERMINOLOGY

- **Database -** One discrete data system; The largest structure we'll be working with
- **Table -** A set of data contained inside a database; Consists of columns and rows

RELATIONAL DATABASE TERMINOLOGY

- **Database -** One discrete data system; The largest structure we'll be working with
- **Table -** A set of data contained inside a database; Consists of columns and rows
- **Record -** One individual set of fields in the database (i.e. a single row)

RELATIONAL DATABASE TERMINOLOGY

- **Database -** One discrete data system; The largest structure we'll be working with
- **Table -** A set of data contained inside a database; Consists of columns and rows
- **Record -** One individual set of fields in the database (i.e. a single row)
- **Key -** A field in a table used to retrieve records (*more later*)

RELATIONAL DATABASE TERMINOLOGY

- **Database -** One discrete data system; The largest structure we'll be working with
- **Table -** A set of data contained inside a database; Consists of columns and rows
- **Record -** One individual set of fields in the database (i.e. a single row)
- **Key -** A field in a table used to retrieve records (*more later*)
- **Schema -** A representation of the underlying structure behind the database; Serves as a template of the database

KEYS AND SQL

- Keys are just a way to retrieve records from our tables
- They should be unique *within* the table
  - Example: You have an `employee` table. What types of keys could you use to identify your records?

KEYS AND SQL

- Keys are just a way to retrieve records from our tables
- They should be unique *within* the table
  - Example: You have an `employee` table. What types of keys could you use to identify your records?
    - **An employee ID**
    - **Driver's license number**
    - **A full legal name (possibly—often causes conflicts!)**

KEYS AND SQL

- There are many different types of keys (we use **primary keys** in CS50):

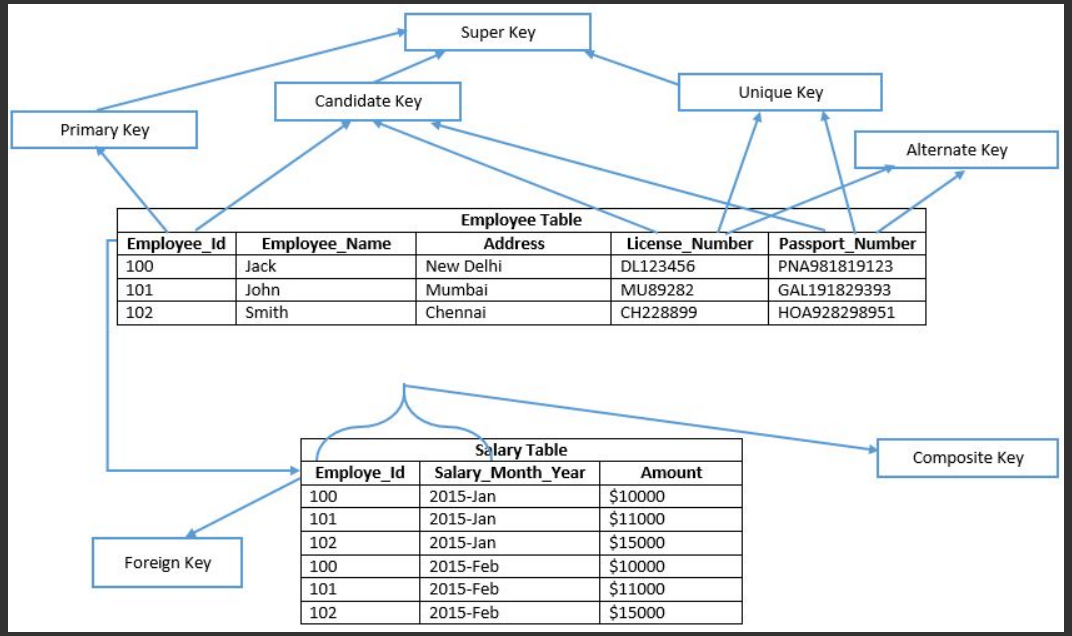

Employee Table

Employee_Id	Employee_Name	Address	License_Number	Passport_Number
100	Jack	New Delhi	DL123456	PNA981819123
101	John	Mumbai	MU89282	GAL191829393
102	Smith	Chennai	CH228899	HOA928298951

Salary Table

Employe_Id	Salary_Month_Year	Amount
100	2015-Jan	$10000
101	2015-Jan	$11000
102	2015-Jan	$15000
100	2015-Feb	$10000
101	2015-Feb	$11000
102	2015-Feb	$15000

STRUCTURED QUERY LANGUAGE (SQL)

- We'll be working with SQL using the **CRUD** paradigm:
  - **Create** new databases, tables, and records (INSERT)
  - **Read** existing records from the database (SELECT)
  - **Update** records and database schema (UPDATE)
  - **Delete** records or database schema (DELETE)

STRUCTURED QUERY LANGUAGE (SQL)

- Really get to know these commands for the Pset:
  - **INSERT**: Inserts values into a table
  - **SELECT**: Selects values from a table
  - **UPDATE**: Updates records and database schema
  - **DELETE**: Deletes records from a table

STRUCTURED QUERY LANGUAGE (SQL)

- SQL is the underlying language that standardizes how we run statements to implement the CRUD paradigm
  - There exists many different **database engines** that allow us to execute SQL
  - We utilize SQLite in CS50, but an extremely popular alternative is MySQL
- We also have access to a tool called `DB Browser` to give us a graphical user interface (GUI) for running different SQL commands

EXAMPLES: THE CONTEXT

users

idnum	username	password	fullname
10	jerry	fus!!!!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza

INSERT

- An `INSERT` query adds information to a table.

```
INSERT INTO
<table>
(<columns>)
VALUES
(<values>)
```

INSERT

- An `INSERT` query adds information to a table.

```
INSERT INTO
users
(username, password, fullname)
VALUES
('newman', 'USMAIL', 'Newman')
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza

INSERT

- When defining the column that ultimately is your primary key, it's usually a good idea for that column to be an integer.

- Moreover, you can configure that column to **autoincrement**, so it will pre-populate that column for you automatically when rows are added, eliminating the risk that you'll accidentally try to insert something with a duplicate value.

INSERT

- An `INSERT` query adds information to a table.

```
INSERT INTO
moms
(username, mother)
VALUES
('kramer', 'Babs Kramer')
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

SELECT

- A `SELECT` query extracts information from a table.

```
SELECT
<columns>
FROM
<table>
WHERE
<predicate>
```

SELECT

- A `SELECT` query extracts information from a table.

```
SELECT
idnum, fullname
FROM
users
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

SELECT

- A `SELECT` query extracts information from a table.

```
SELECT
password
FROM
users
WHERE
idnum < 12
```

users

idnum	username	password	fullname
10	jerry	fus!!!!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

SELECT

- A `SELECT` query extracts information from a table.

```
SELECT
*
FROM
moms
WHERE
username = 'jerry'
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

JOIN

- Databases empower us to organize information into tables efficiently.

- We don’t always need to store every possible relevant piece of information in the same table, but rather we can use relationships across tables to connect all the pieces of data we need.

- Let’s imagine we need to get a user’s full name (from the *users* table) and their mother’s name (from the *moms* table).

JOIN

- A `SELECT (JOIN)` query extracts information from multiple tables.

```
SELECT
<columns>
FROM
<table1>
JOIN
<table2>
ON
<predicate>
```

JOIN

- A `SELECT (JOIN)` query extracts information from multiple tables.

```
SELECT
users.fullname, moms.mother
FROM
users
JOIN
moms
ON
users.username = moms.username
```

JOIN

- A `SELECT (JOIN)` query extracts information from multiple tables.

```
SELECT
users.fullname, moms.mother
FROM
users
JOIN
moms
ON
users.username = moms.username
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

users & moms

users.idnum	users.username moms.username	users.password	users.fullname	moms.mother
10	jerry	fus!ll!	Jerry Seinfeld	Helen Seinfeld
11	gcostanza	b0sc0	George Costanza	Estelle Costanza

users & moms

users.idnum	users.username moms.username	users.password	users.fullname	moms.mother
10	jerry	fus!ll!	Jerry Seinfeld	Helen Seinfeld
11	gcostanza	b0sc0	George Costanza	Estelle Costanza

UPDATE

- An `UPDATE` query modifies information in a table.

```
UPDATE
<table>
SET
<column> = <value>
WHERE
<predicate>
```

UPDATE

- An `UPDATE` query modifies information in a table.

```
UPDATE
users
SET
password = 'yadayada'
WHERE
idnum = 10
```

users

idnum	username	password	fullname
10	jerry	yadayada	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

DELETE

- A `DELETE` query removes information from a table.

```
DELETE FROM
<table>
WHERE
<predicate>
```

DELETE

- A `DELETE` query removes information from a table.

```
DELETE FROM
users
WHERE
username = 'newman'
```

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza
12	newman	USMAIL	Newman

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

users

idnum	username	password	fullname
10	jerry	fus!ll!	Jerry Seinfeld
11	gcostanza	b0sc0	George Costanza

moms

username	mother
jerry	Helen Seinfeld
gcostanza	Estelle Costanza
kramer	Babs Kramer

STRUCTURED QUERY LANGUAGE (SQL)

SQL Cheat Sheet

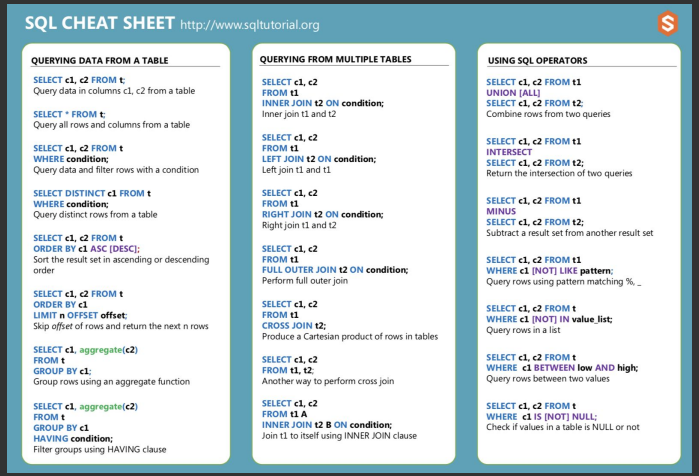

SQL CHEAT SHEET http://www.sqltutorial.org

QUERYING DATA FROM A TABLE

SELECT c1, c2 FROM t;
Query data in columns c1, c2 from a table

SELECT * FROM t;
Query all rows and columns from a table

SELECT c1, c2 FROM t
WHERE condition;
Query data and filter rows with a condition

SELECT DISTINCT c1 FROM t
WHERE condition;
Query distinct rows from a table

SELECT c1, c2 FROM t
ORDER BY c1 ASC [DESC];
Sort the result set in ascending or descending order

SELECT c1, c2 FROM t
ORDER BY c1
LIMIT n OFFSET offset;
Skip *offset* of rows and return the next n rows

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1;
Group rows using an aggregate function

SELECT c1, aggregate(c2)
FROM t
GROUP BY c1
HAVING condition;
Filter groups using HAVING clause

QUERYING FROM MULTIPLE TABLES

SELECT c1, c2
FROM t1
INNER JOIN t2 ON condition;
Inner join t1 and t2

SELECT c1, c2
FROM t1
LEFT JOIN t2 ON condition;
Left join t1 and t1

SELECT c1, c2
FROM t1
RIGHT JOIN t2 ON condition;
Right join t1 and t2

SELECT c1, c2
FROM t1
FULL OUTER JOIN t2 ON condition;
Perform full outer join

SELECT c1, c2
FROM t1
CROSS JOIN t2;
Produce a Cartesian product of rows in tables

SELECT c1, c2
FROM t1, t2;
Another way to perform cross join

SELECT c1, c2
FROM t1 A
INNER JOIN t2 B ON condition;
Join t1 to itself using INNER JOIN clause

USING SQL OPERATORS

SELECT c1, c2 FROM t1
UNION [ALL]
SELECT c1, c2 FROM t2;
Combine rows from two queries

SELECT c1, c2 FROM t1
INTERSECT
SELECT c1, c2 FROM t2;
Return the intersection of two queries

SELECT c1, c2 FROM t1
MINUS
SELECT c1, c2 FROM t2;
Subtract a result set from another result set

SELECT c1, c2 FROM t1
WHERE c1 [NOT] LIKE pattern;
Query rows using pattern matching %, _

SELECT c1, c2 FROM t
WHERE c1 [NOT] IN value_list;
Query rows in a list

SELECT c1, c2 FROM t
WHERE c1 BETWEEN low AND high;
Query rows between two values

SELECT c1, c2 FROM t
WHERE c1 IS [NOT] NULL;
Check if values in a table is NULL or not

sqlite3 Cheat Sheet

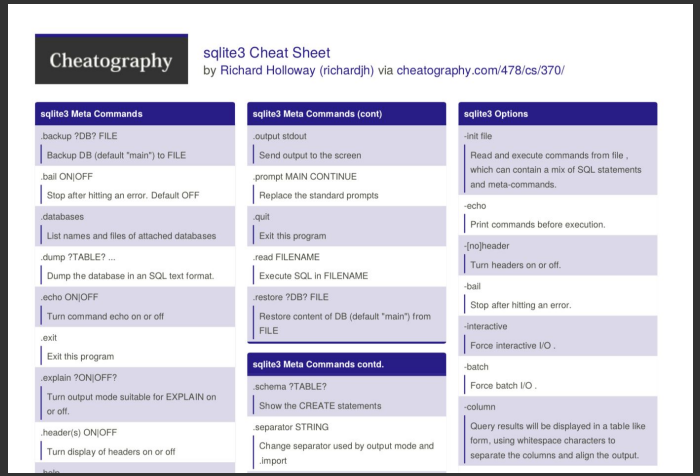

Cheatography

sqlite3 Cheat Sheet
by Richard Holloway (richardjh) via cheatography.com/478/cs/370/

sqlite3 Meta Commands

.backup ?DB? FILE
Backup DB (default "main") to FILE

.bail ON|OFF
Stop after hitting an error. Default OFF

.databases
List names and files of attached databases

.dump ?TABLE? ...
Dump the database in an SQL text format.

.echo ON|OFF
Turn command echo on or off

.exit
Exit this program

.explain ?ON|OFF?
Turn output mode suitable for EXPLAIN on or off.

.header(s) ON|OFF
Turn display of headers on or off

sqlite3 Meta Commands (cont)

.output stdout
Send output to the screen

.prompt MAIN CONTINUE
Replace the standard prompts

.quit
Exit this program

.read FILENAME
Execute SQL in FILENAME

.restore ?DB? FILE
Restore content of DB (default "main") from FILE

sqlite3 Meta Commands contd.

.schema ?TABLE?
Show the CREATE statements

.separator STRING
Change separator used by output mode and .import

sqlite3 Options

-init file
Read and execute commands from file , which can contain a mix of SQL statements and meta-commands.

-echo
Print commands before execution.

-[no]header
Turn headers on or off.

-bail
Stop after hitting an error.

-interactive
Force interactive I/O .

-batch
Force batch I/O .

-column
Query results will be displayed in a table like form, using whitespace characters to separate the columns and align the output.

Demo

Creating a Database

Now Let’s Do Some Exercises!

http://bit.ly/2oojaps

Exercise Solutions

https://github.com/rw5614/cs50sectiontf/blob/master/solution.sql

REFERENCE SHEETS

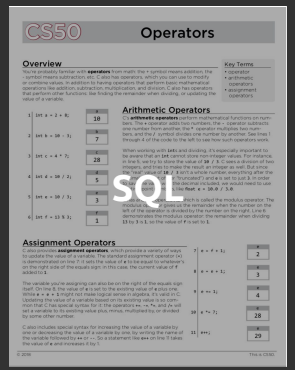

https://www.dropbox.com/sh/5y662ey1hc4sde4/AACjqHN3NtSKk4ShsRDFd_Sja?dl=0&m=&preview=SQL.pdf

THE WEEK AHEAD

Final Live Lecture: Fri 11/1

Pset 7: Sun 11/3 @ 11:59pm

- `Movies`
- `Houses`

Test:

- Review Session: Mon 11/4 @ 3-5pm, Sanders (Also Posted Online)
- Released Mon 11/4 @ 7:30pm
- Due Fri Nov 8 @ 11:59pm

TEST REVIEW TIPS

TEST (11/4 - 11/8)

- Covers Week 0 Through 7
- Open book, *Conceptually-Focused*
- Released via Course Website
- Submitted via submit50
- Expect to spend several hours (not days!)
- Preparation Guidelines: https://cs50.harvard.edu/college/test/

TIPS FOR THE TEST

- Review lecture notes, these slides, reference sheets, and other conceptually-focused materials
- We've spent a lot of time discussing context and *why* in section—Apply that mindset to the test
  - Ask: Why do we use a certain algorithm? What role does this play in my code? Is this syntax or convention?
- Work out the Test problems on paper and type them up after the fact
  - Remember our whiteboarding strategy
  - If you're trying to come up with an algorithm, run through examples, **edge cases**, etc.

ONE MORE THING

It's been a privilege to be your Teaching Fellow these weeks.

Here to be a resource for you all now and *after* CS50:

- raymond_wang@college.harvard.edu
- (617) 949-6919

THANKS FOR BEING A FANTASTIC SECTION!